

# INT354

# Machine Learning Foundations

---

Machine Learning Classifiers using  
Scikit-learn

# Introduction to scikit-learn

---



- **Machine Learning in Python**
- **Simple and Effective tools for data mining and data Analysis**
- **Built on Numpy, Scipy, Matplotlib**
- **Open source, commercially usable - BSD license**
- **Accessible to everybody, reusable in various contexts**

# How to Install?

---

Python Idle

- **pip install scikit-learn**

Anconda

- **conda install scikit-learn**

Jupyter Note Book

- **!pip install scikit-learn**



# No Free Lunch Theorem

---

- **No Classifier works best across all possible scenarios.**



# Choosing a classification Algorithm?

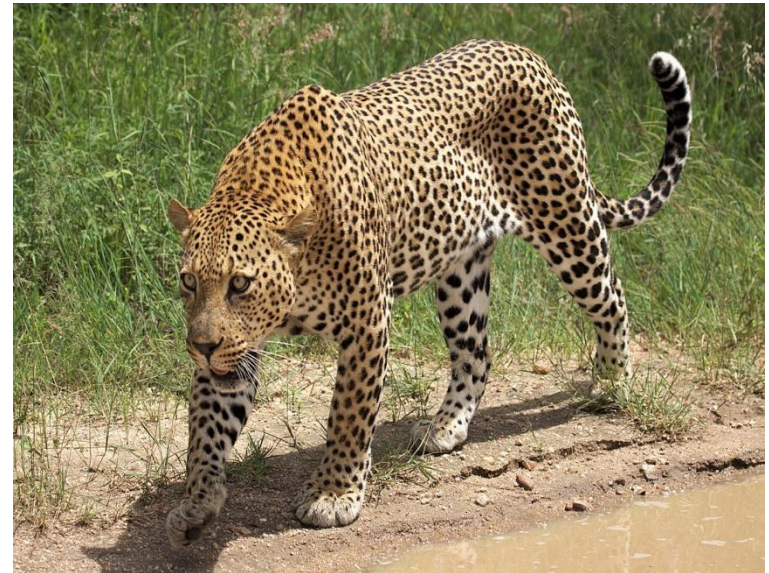
---

## Parameters for choosing a classification algorithm

- **Performance may depends on**
  - **Number of features**
  - **Number of samples**
  - **The amount of noise in dataset**
  - **Classes are linearly separable or NOT.**

# Selection of Features

Features to classify dissimilar Object



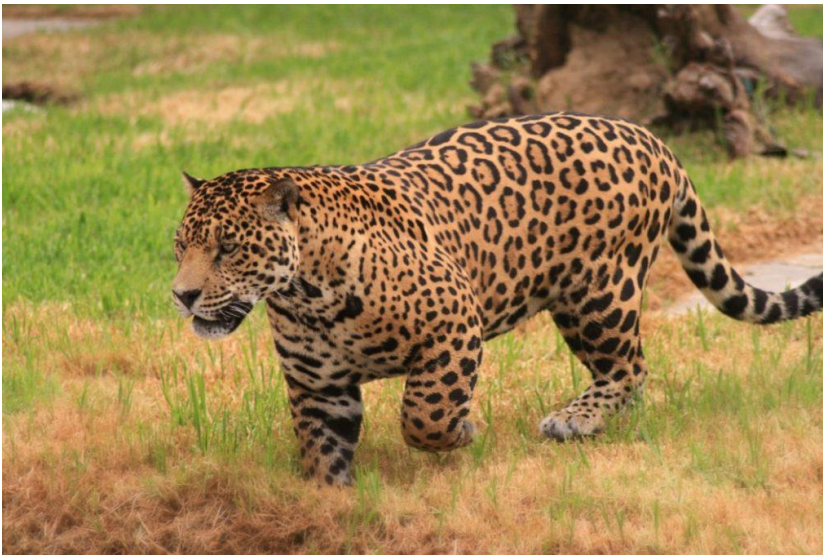
Tiger Vs Leopard

Easy??



# Selection of Features

Features to classify similar Object



Jaguar Vs Leopard

??

# Training a machine learning algorithm

---

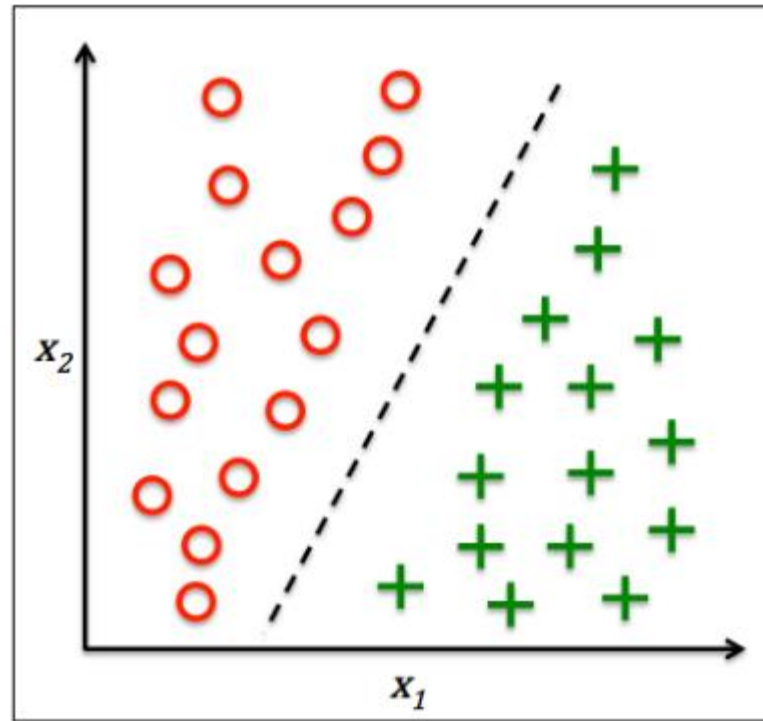
Training a machine learning algorithm can be summarized in 5 main steps.

- 1. Selection of features**
- 2. Choosing a performance Metrics**
- 3. Choosing a classifier and Optimization Algorithm.**
- 4. Evaluating the performance of model**
- 5. Tuning the algorithm.**



# Perceptron

- 2 class classifier
- Linear separable patterns



# Training Perceptron using scikit-learn

---

- **Major Steps involved**
  1. Load dataset
  2. Divide the data in train and test
  3. Feature Scaling
  4. Train Perceptron model
  5. Check accuracy
  6. Plot Decision Boundary

# Available Datasets in scikit

---

- `load_boston()`
- `load_iris()`
- `load_diabetes()`
- `load_digits()`
- `load_linnerud()`
- `load_wine()`
- `load_breast_cancer()`

# Load Dataset

---

```
from sklearn import datasets  
Iris=datasets.load_iris()  
x=iris.data  
y=iris.target
```

# Split data in train and test

---

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test =
train_test_split(x,y,test_size=0.2,random_state=0)
print(np.shape(x))
print(np.shape(x_train))
print(np.shape(x_test))
```

# Feature Scaling

---

```
from sklearn.preprocessing import StandardScaler  
sc=StandardScaler()  
sc.fit(x_train)  
x_train_std=sc.transform(x_train)  
x_test_std=sc.transform(x_test)
```

## Note:

- `fit()` method estimates the sample mean( $\mu$ ) and standard deviation ( $\sigma$ ) for each feature dimension from the training data.
- `transform()` method standardized the training data and test data using the estimated parameters.

# Train Perceptron Model

---

```
from sklearn.linear_model import Perceptron  
ppn=Perceptron(eta0=0.1, random_state=0)  
ppn.fit(x_train_std, y_train)
```

## Note:

- If learning rate is too large, then algorithm will overshoot the global cost minimum.
- If learning rate is too small, then algorithm requires more epochs until convergence.



# Prediction on new data

---

```
y_pred=ppn.predict(x_test_std)  
print('Misclassified Samples: %d' %(y_test != y_pred).sum())
```

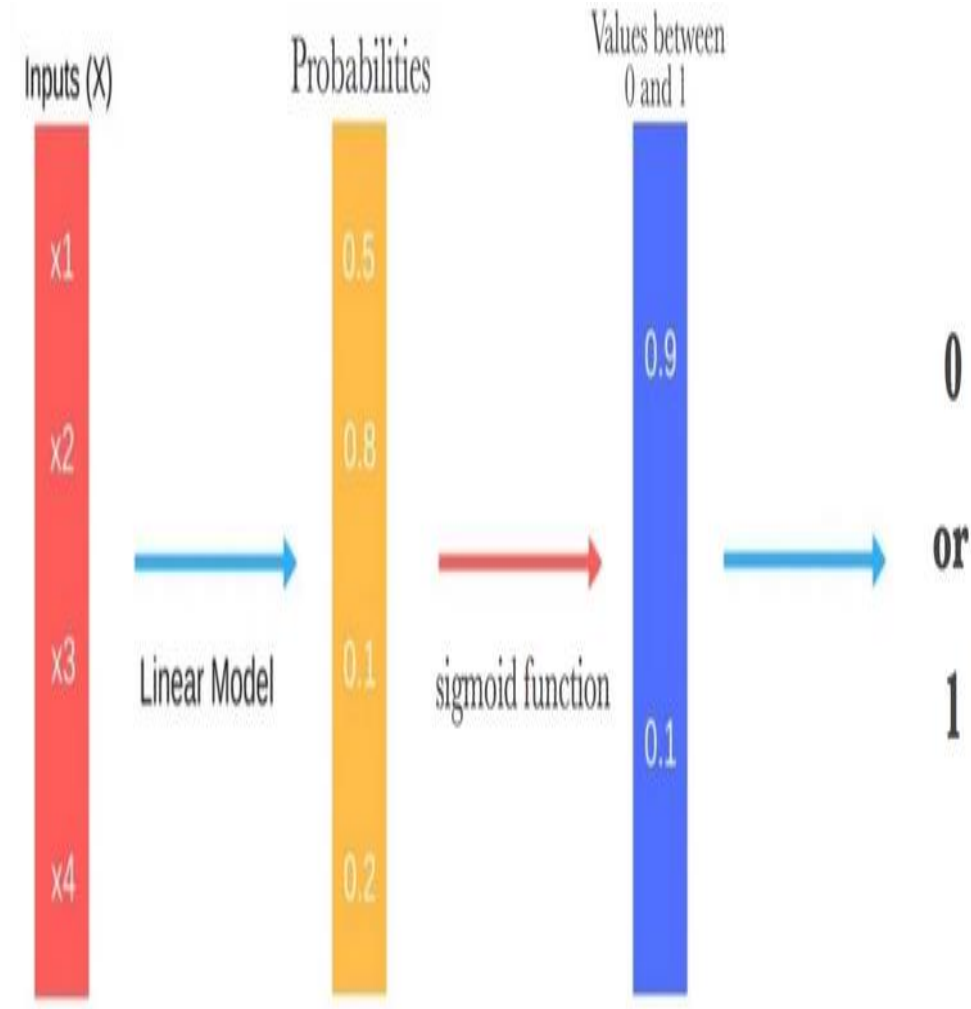
# Compute Accuracy

---

```
from sklearn.metrics import accuracy_score  
print('Accuracy: %f '%accuracy_score(y_test,y_pred))
```

# Logistic Regression and Conditional Probabilities

- **Measures the relationship between a target and one or more independent variables by plotting the dependent variable's probability score.**
- **For example:**
  - **To predict whether an email is spam or not.**
  - **Whether the tumor is malignant or not.**



# Odds Ratio

---

$$\text{Odds\_Ratio} = \frac{P}{1-P}$$

Where P stands for the Probability of positive event.

$$\text{logit}(p) = \log \frac{p}{(1-p)}$$

$$\text{logit}(p(y = 1|x)) = wTx$$

To predict the probability whether a sample belongs to certain class, take the inverse of logit function.

$$\phi(z) = \frac{1}{1 + e^{-z}}$$

Where z is net input.

$$z = wTx = w_0 + w_1x_1 + \dots + w_m x_m$$

# Activation Function in Logistic Regression


---

## Sigmoid Function

$$f(x) = \frac{1}{1 + e^{-x}}$$

# Decision Boundary

- Set threshold to predict the class of sample.
- Can be linear or non-linear.
- Increase polynomial order to get complex decision boundary.



# Cost Function

---

$$j(\mathbf{w}) = \sum_{i=1}^n -\log(\phi(\mathbf{z}^{(i)})) - (1 - y^{(i)})\log(1 - \phi(\mathbf{z}^{(i)}))$$

$$j(\phi(\mathbf{z}), y; \mathbf{w}) = \begin{cases} -\log(\phi(\mathbf{z})) & \text{if } y = 1 \\ -\log(1 - \phi(\mathbf{z})) & \text{if } y = 0 \end{cases}$$



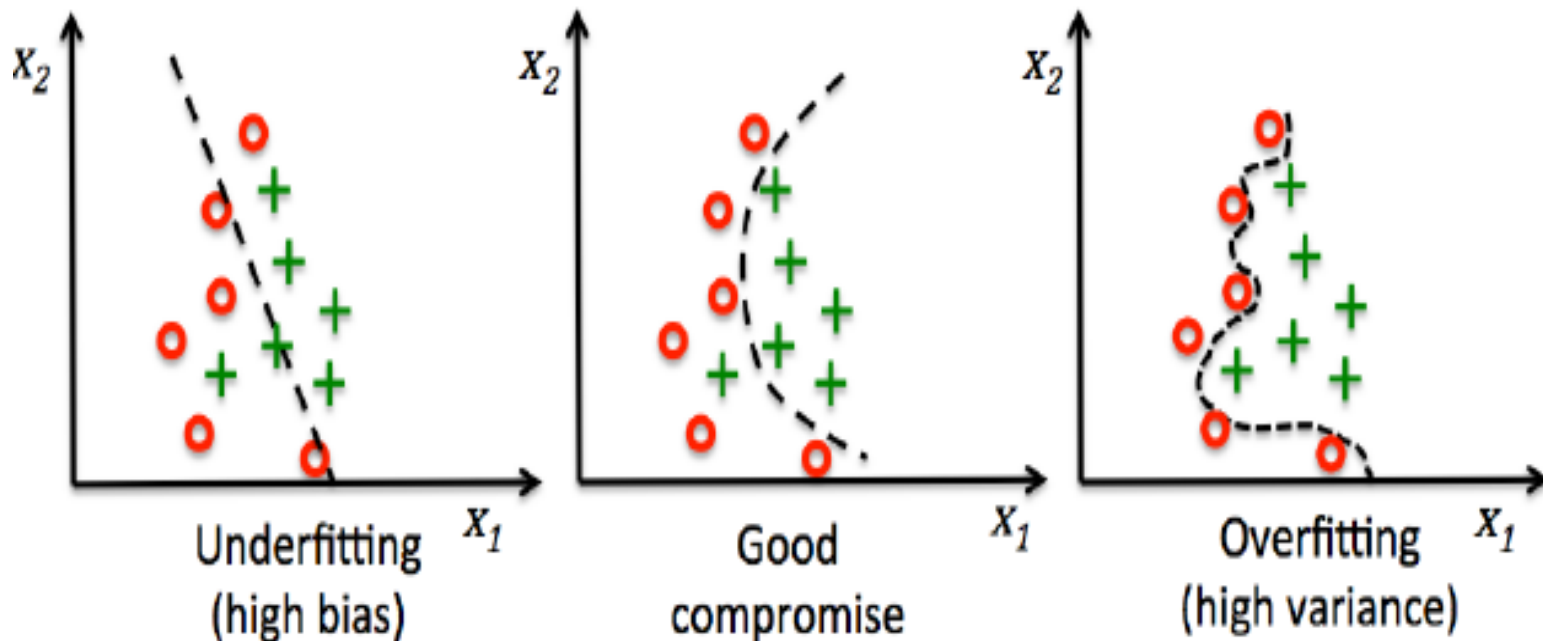
# Weight Updating in Logistic Regression

---

$$w_j = w_j + \eta \sum_{i=1}^n (y^{(i)} - \phi(z^{(i)}))x^{(i)}$$
$$w := w + \Delta w$$

$$\Delta w_j = -\eta \frac{\partial J}{\partial w_j} = \eta \sum_{i=1}^n (y^{(i)} - \phi(z^{(i)}))x^{(i)}$$

# Over-fitting and Under-fitting



## Note:

- Bias measures how far off the predictions from the correct values.
- Variance measures the consistency of the model prediction for a particular sample instance

# Regularization

---

- Handle collinearity
- Filters out noise from data.
- Prevent over-fitting.
- Introduces additional information (bias) to penalize parameter weights.

# L2 Regularization

$$\frac{\lambda}{2} ||w||^2 = \frac{\lambda}{2} \sum_{j=1}^m w_j^2$$

Here,  $\lambda$  is regularization parameter.

To apply regularization, add regularization term to the cost function.

$$J(w) = \left[ \sum_{i=1}^n -\log(\phi(z^{(i)})) - (1 - y^{(i)}) \log(1 - \phi(z^{(i)})) \right] + \frac{\lambda}{2} ||w||^2$$

By increasing the value of  $\lambda$ , regularization strength can be increased.

# Training logistic regression model with scikit-learn

---

```
from sklearn.linear_model import LogisticRegression  
lr=LogisticRegression(C=1000, random_state=0)  
lr.fit(x_train_std, y_train)
```

# Advantages of logistic regression

---

- **Doesn't require input features to be scaled.**
- **Doesn't require any tuning.**
- **Easy to implement.**

# Disadvantages of logistic regression



---

- **Can't solve non-linear problems.**



# Support Vector Machine

- **Objective is to maximize the margin.**



# Maximum margin

$$w_0 + w^T x_{pos} = 1 \quad (1)$$

$$w_0 + w^T x_{neg} = -1 \quad (2)$$

Subtract eq. 1 and 2

$$w^T (x_{pos} - x_{neg}) = 2$$

Normalize this by the length of vector  $w$ ,

$$||w|| = \sqrt{\sum_{j=1}^m w_j^2}$$

So, equation becomes:

$$\frac{w^T (x_{pos} - x_{neg})}{||w||} = \frac{2}{||w||}$$

Left side of equation is called **margin**.

# Non-linear Separable Data

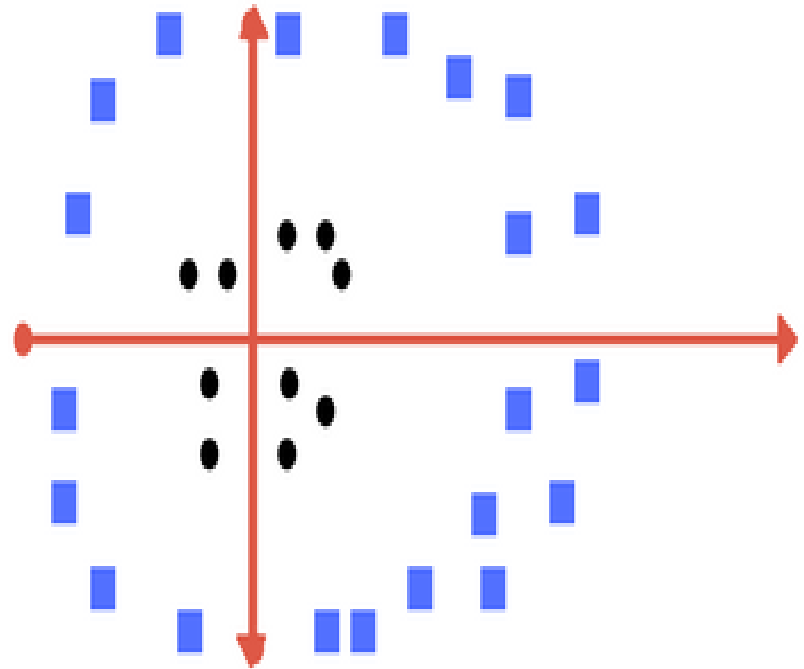
---

# Non-linear separable case

- Slack variables are added for nonlinearly separable data to allow convergence.
- Positive values slack variables are added to the linear constraints:

$$w^T x^{(i)} \geq 1 \text{ if } y^{(i)} = 1 - \xi^{(i)}$$

$$w^T x^{(i)} < -1 \text{ if } y^{(i)} = 1 + \xi^{(i)}$$

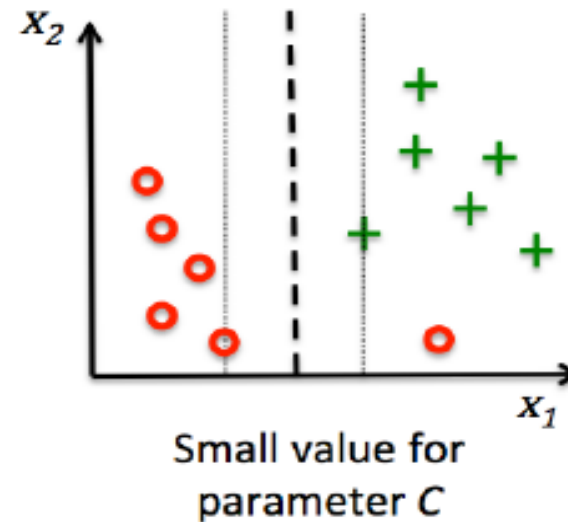
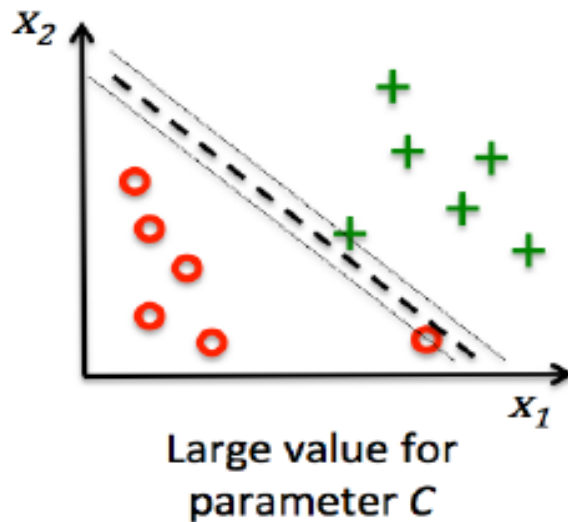


# Objective Function

Minimize

$$\frac{1}{2} ||\mathbf{w}'||^2 + C(\sum_i \xi^{(i)})$$

- C controls the penalty for misclassification.
- C is used to control the width of margin.



# Support Vector Machine (Kernels)

---

- **Linear Kernel SVM**
- **Polynomial Kernel SVM**
- **Radial Kernel SVM**

# Linear Kernel SVM

---

- Dot-product

$$K(x, xi) = \text{sum}(x * xi)$$

- Defines similarity or distance measure between new data and the support vectors.



# Polynomial Kernel SVM

---

$$K(x, xi) = 1 + \text{sum}(x * xi)^d$$

- **d is the degree of polynomial.**

# Radial Kernel SVM

---

$$K(x, x_i) = \exp(-\gamma \sum (x - x_i)^2)$$

- Good default value for gamma is 0.1
- Gamma lies in range of 0 to 1.

# SVM Advantages

---

- Works well with unstructured data like text and images.
- The kernel trick is real strength of SVM.
- The risk of over-fitting is less in SVM.
- It scales relatively well to high dimensional data.

# SVM Disadvantages

---

- Choosing a “good” kernel function is not easy.
- Long training time for large datasets.
- Difficult to understand and interpret the final model.



**COMING UP**

---